

Fog Project Documentation

Edge-Cloud application simulating household power consumption
with robust disconnection handling

Marvin Steinke

Niklas Fomin

July 11, 2023

This project simulates a distributed power system where power data from different households is streamed in real-time. The data is processed and averaged over a 5-second window. The simulation is coordinated by Mosaik, a co-simulation framework, and uses Apache Kafka for handling real-time data. The averaged data gets queued into a SQLite database and arranged into a schema that allows for the tracking of each inserted tuple. The edge server manages the messaging of the queued data by continuously checking connectivity to the cloud server and buffering the generated tuples into the sockets. The cloud server caches each received tuple and generates the corresponding postal code for the household, and sends this back to the edge server. The edge server eventually stores this in its database and marks the tuple as acknowledged. The initial start of the cloud server triggers a backend running on the same hosts that fetches key-value pairs out of the cache and visualizes them on charts, accessible on the web. A schema for the project is depicted in Figure 1.

1 Edge Component

- **data:** Based on the original data from “Gem House Opendata: German Electricity Consumption in Many Households Over Three Years 2018-2020 (Fresh Energy)” [?]. The data is reduced to a one hour window and made compatible with the simulation environment by adding a title and converting the timestamp from `YYYY-MM-DD HH:mm:ss.SSS` to `YYYY-MM-DD HH:mm:ss`.
- **household.py:** Contains the Household simulator that generates power data based on the provided CSV file.
- **kafka_adapter.py:** Contains the KafkaAdapter and KafkaAdapterModel classes, responsible for connecting the Mosaik simulation with the Kafka data stream.
- **collector.py:** Contains the Collector simulator that collects all power data.
- **main.py:** Contains the main function for starting the simulation.

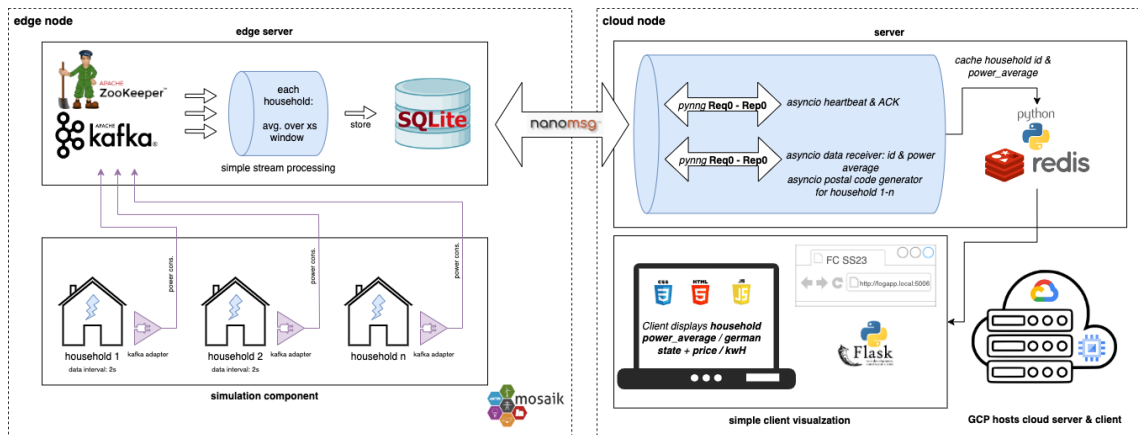


Figure 1: Project Schema

- **edge_server.py**: Contains the `EdgeServer` class for reading, writing, and inserting the produced data from the Kafka topics into SQLite. Additionally, three concurrent threads are used to continuously check connectivity, publish data read from the Kafka topics and the SQLite database, and subscribe to messages from the cloud server. For messaging, we use `Pynng`, which is a Python wrapper of the lightweight, high-performance messaging library `nano message`. Initially, it does not provide “reliable” features such as message buffering or connection retries. It is considered to be based on `ZeroMQ` but aims to simplify its usage. With concerns to reliability the edge server uses the database to track if a message was sent or not. Initially the flag is set to 0. If sent it turn 1. Only if the server received it and replied with the according postal code of the received ID the message is considered acknowledged and the flag is set to 2 and eventually the postal code is inserted into the db. Whenever the connection to the cloud node is established the data sender thread wakes up and both fetches both new data with sent flag 0 or data that might be lost with the flag 1. This mechanism contributes to the robust disconnection handling. whenever the cloud-node crashes or the connection is disrupted the edge-server will buffer the queued messages to be sent. When the edge-server crashes the cloud node logs accordingly and continues retrying to acknowledge heartbeats from the edge side. As soon as the edge-server is back online it benefits of the persistent data-storage and continues sending data beginning from the last sent message ID. Again, in case the cloud-node sent the acknowledgement, but the edge-node crashed, as soon as it starts up again it will check the database and resent the unsent or unacknowledged messages to the cloud-node.
- **db_operations.py**: The database contains operations for the lightweight, file-based SQLite database used by the `EdgeServer`. It creates a schema, inserts tuples received from the Kafka producer thread, and fetches newly inserted and lost data by tracking and checking the `sent` flag. Furthermore, it handles the insertion of messages received from the cloud server.
- **reliable messaging**: In terms of reliability, the edge server employs a database to track the status of messages. To ensure robust handling of disconnections, the data sender thread wakes up whenever the connection to the cloud node is established. The edge server retrieves two types of data: newly queued data with the flag set to 0, indicating it hasn’t been sent yet, and data that may have been previously sent but not acknowledged, with the flag set to 1. It then buffers this data into the socket for further processing. This mechanism helps in managing potential disruptions and maintaining a reliable connection. If the cloud node crashes or the connection is disrupted, the edge server buffers the queued messages to be sent. In the event of an edge server crash, the cloud node logs the occurrence and continues attempting to acknowledge heartbeats from the edge side. Once the edge server comes back online, it benefits from the persistent data storage and resumes sending data starting from the last sent message ID. Additionally, if the cloud node successfully sends an acknowledgement but the edge node crashes, the edge node checks the database upon restarting and resends any unsent or unacknowledged messages to the cloud node. This ensures that no data is lost in case of an edge server failure.

2 Cloud Node

The cloud node runs the `cloud_server.py` file to interact asynchronously with the edge component through two `pynng` sockets, utilizing the `Req/Rep` message pattern. The cloud node awaits data from the Kafka server and caches the received data into a Redis cache. For each received message it generates a postal code and sends it back to the edge server and additionally triggers the `data_fetcher.py` script in the backend to fetch data from the cache. This process provides a REST API for an ajax engine running `main.js` in the frontend. The cloud server responds to continuous heartbeat messages to ensure necessary reliability and handle disconnections in the event of network failures or application crashes. Both the cloud server and the web client can be operated by the startup `script.sh` integrated in the terraform files.

- **cloud_server.py**: Contains asynchronous (implemented with `asyncio` & `pynng`) functions that reply to heartbeat messages by the edge-server, receive the power data and reply with the generated postal codes back to the edge-server. The received data gets cached to redis-server instance on the same host.
- **data_fetcher.py**: Contains the flask backend script to fetch data out of the redis-cache, associate one of sixteen german states with the id of the received data tuple and providing a

REST API for the frontend to present the data values on a webpage.

- `main.js + index.html + main.css`: Contains a simple frontend logic to access the key-value-pairs using ajax and implements a bar-chart of the `Chart.js` library to visualize the average power-consumption per german state in annual kw/h + the associated price based on an assumption.